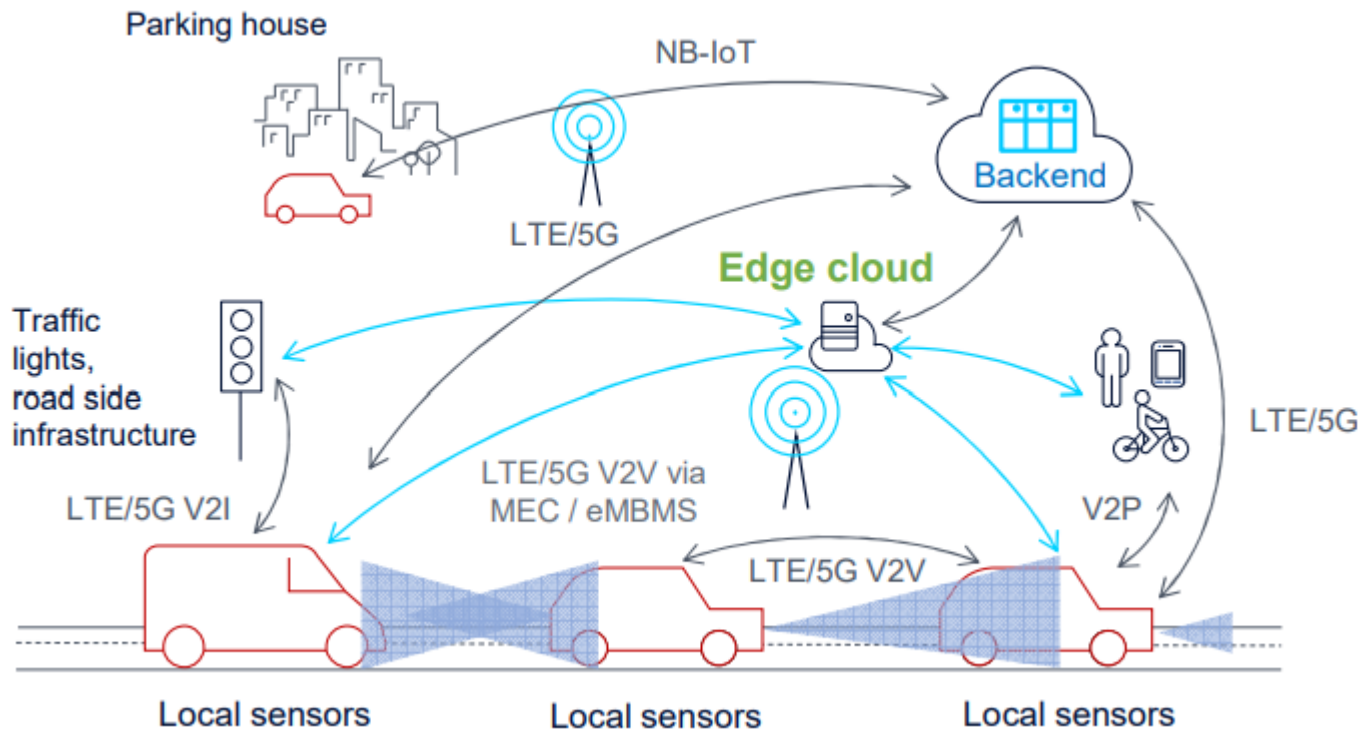# Self-organized systems: Data, learning and decisions
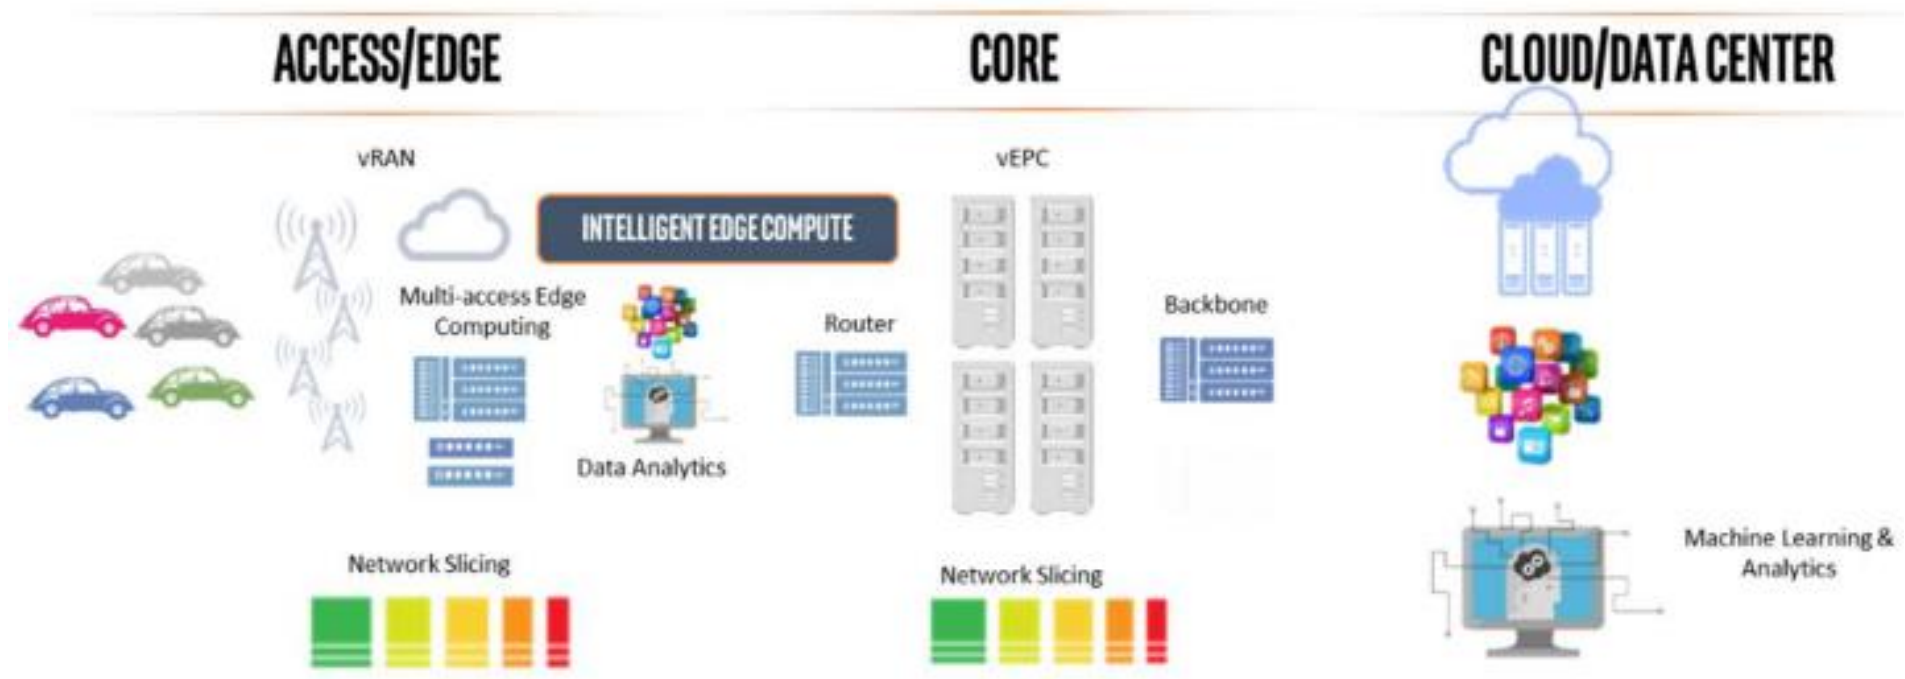
Mestrado em Engenharia de Computadores e Telemática

2022/2023

Use cases and data

# Where to process the data?

Data Analytics

- Processing data
- Get some decision with the data
- Network decisions with network and services data
  - Give more bandwidth?
  - Assign some special queue to reduce the delay?
- User decisions or element decisions
  - Obstacle in place?
  - Robot kicks the ball to the right?
- Network decisions with users' data
  - Predict handovers with location and velocity
  - Move the CDN content to the users' most near access point
  - Ambulance is on the way with network requirements
- User decisions with network data
  - Chose a path with great connectivity for gaming or video
 - Chose a place for remote augmented and virtual reality

Machine Learning

- A pragmatic definition:
  - Collection of algorithms and statistical models (methods) for machines to carry out automated tasks *based on the observation of inputs and/or outputs of a process*
- The goal of Machine Learning is to produce an estimation or a classification given a set of input values.
- We often distinguish:
  - ML method: the mechanism to train a model (neural network, support vector machine, etc.)
  - ML model: an instance of the method trained to replicate the behavior of the target process

Types of Learning

- Supervised: model is trained with a dataset of the target process
    - When training for a classification task, the historical dataset should contain the Ground Truth - the actual class of a given sample
- Unsupervised: classification or regression does not depend on prior knowledge

Ground Truth necessary for training

Historical Dataset		
Feature 1	Feature 2	Class
A	1	X
B	1	X
C	2	Y

Unsupervised (classes are created by, e.g., finding clusters of similar data points)

Supervised (classification depends on historical inputs)

Supervised Learning

- Training stage

A dataset with input data and corresponding output

The input data is pre-processed to identify and/or extract relevant features

The feature data is input to the ML method, typically one feature set (e.g., mean, median, std. dev.)

      Sometimes a blind set of features is produced, and then only the most relevant are selected (e.g., decision trees)

For each input set, the method produces an estimate likely to have an error.

The method compares the estimate with the actual process output (the Ground Truth), and updates the model's internal processes to improve the accuracy of the estimates.

The process is repeated until performance of the method is within acceptable bounds.

- Inference stage

The trained model is deployed in its target setting.

Given inputs, it can produce estimates of the process output.

However, the method no longer has access to the ground truth,

and it thus enable of further learning.

Training stage

Set 1: <input1, input2, inputN> → <feat1, feat2, featM>
Set 2: <input1, input2, inputN> → <feat1, feat2, featM>
Set 3: <input1, input2, inputN> → <feat1, feat2, featM>
...
Set S: <input1, input2, inputN> → <feat1, feat2, featM>

Set 1: <input1, input2, inputN; output>
Set 2: <input1, input2, inputN; output>
Set 3: <input1, input2, inputN; output>
...
Set S: <input1, input2, inputN; output>

Raw dataset <Inputs, Outputs> → Extract features → *Features* → Train model → *Output* → Evaluate Accuracy

Inference stage

Run-time Data → *Raw data* → Extract features → *Features* → Trained model → *Output*

Ground truth
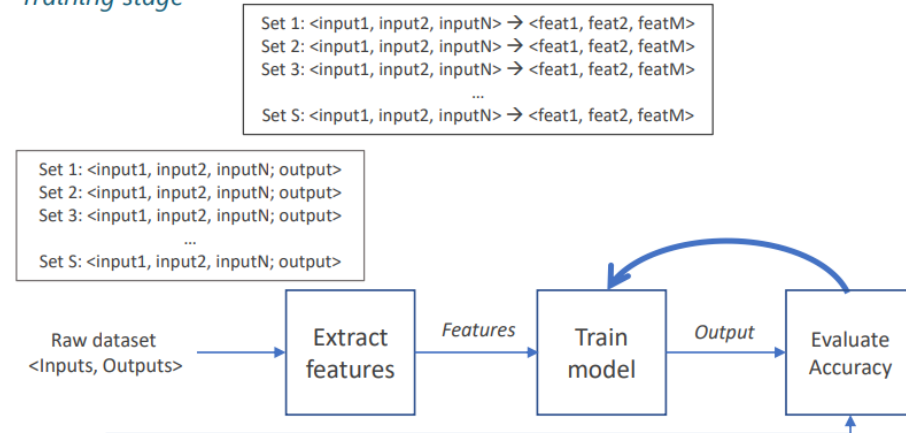
Neural networks

- One of the most successful methods of ML
- Building blocks: Perceptron and Synapse
- Perceptron: typically a function that maps the entire natural range into a bounded interval ([0,1] or [-1,1])
  - Example: Logistic sigmoid function, ReLU function, tanh, softmax, etc
- Synapses: connections from perceptrons of layer (n-1)th to perceptrons of layer nth, each applying a weight to the transmitted value
- Training Neural Networks is mostly about finding the weights of those synapses

Logistic sigmoid function

$$\sigma = \frac{1}{1+e^{-v}}$$
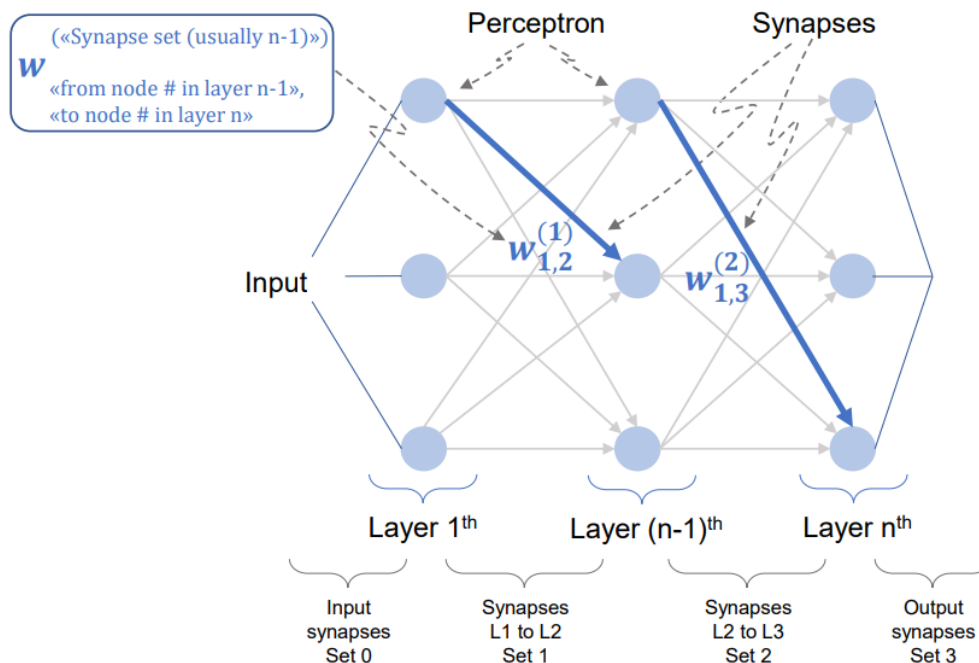
Reinforcement Learning

■ Type of machine learning technique that enables an agent to learn in an interactive environment by trial and error using feedback from its own actions and experiences

Types of Machine Learning

1. **Environment** — Physical world in which the agent operates
2. **State** — Current situation of the agent
3. **Reward** — Feedback from the environment
4. **Policy** — Method to map agent's state to actions
5. **Value** — Future reward that an agent would receive by taking an action in a particular state

$$Q(s_t, a_t) \leftarrow \underbrace{(1-\alpha) \cdot \underbrace{Q(s_t, a_t)}_{\text{old value}}}_{} + \underbrace{\alpha}_{\text{learning rate}} \cdot \left(\overbrace{\underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}}}^{\text{learned value}} \right)$$

■ **Q-learning:** updates Q values which denotes value of performing action a in state s. The following value update rule is the core of the Q-learning algorithm.

Reward example: best path with resources: path bandwidth / path length

Learning rate and discount factor:]0 1[

$$reward = \frac{\frac{100.0}{|P|} \times \sum_{l}^{l \in P} R_l}{E_{pl} \times K_{spl}}$$

Computer games (pacman example), robot automation (RL is used to enable the robot to create an efficient adaptive control system for itself which learns from its own experience and behavior**), …**

Unsupervised learning: K-means

- Centroid: non-data point that indicates center of cluster as identified by K-means

- Operation:
  1. Deploy N centroids randomly (N proportional to number of expected classes)
  2. Assign randomly data points to classes
  3. Repeat iteratively
     1. Compute center of gravity of each class;
     2. Centroid is repositioned in that center of gravity
     3. Update boundary
  4. Stop when updates become negligible

Centroids
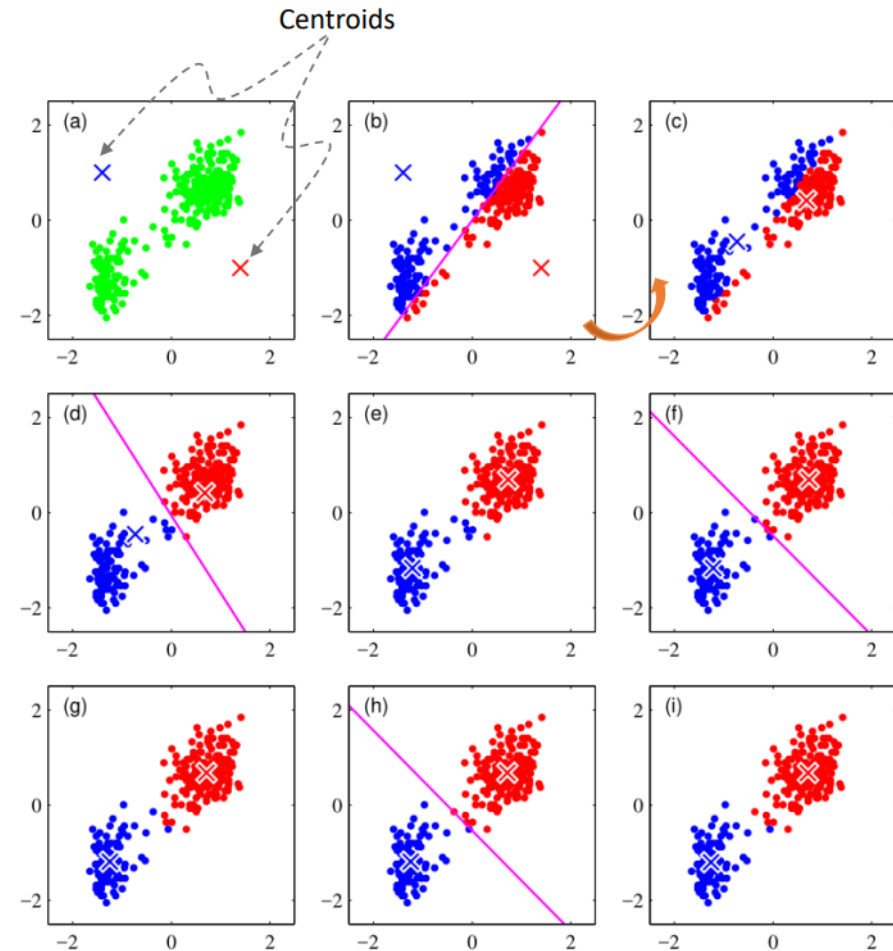
Learning: centralized

Traditional centralized learning — ML runs in the cloud, gathering info from all connected devices and sending back a model.

- The model can generalize based on data from a group of devices and thus instantly work with other compatible devices
- Data can explain all variations in the devices and their environment
- Connectivity - data must be transmitted over a stable connection
- Bandwidth – e.g. a new electrical substation could generate 5 GB/s
- Latency - real-time applications, e.g. automation, requires very low latency
- Privacy - sensitive operational data must remain on site

Learning: de-centralized
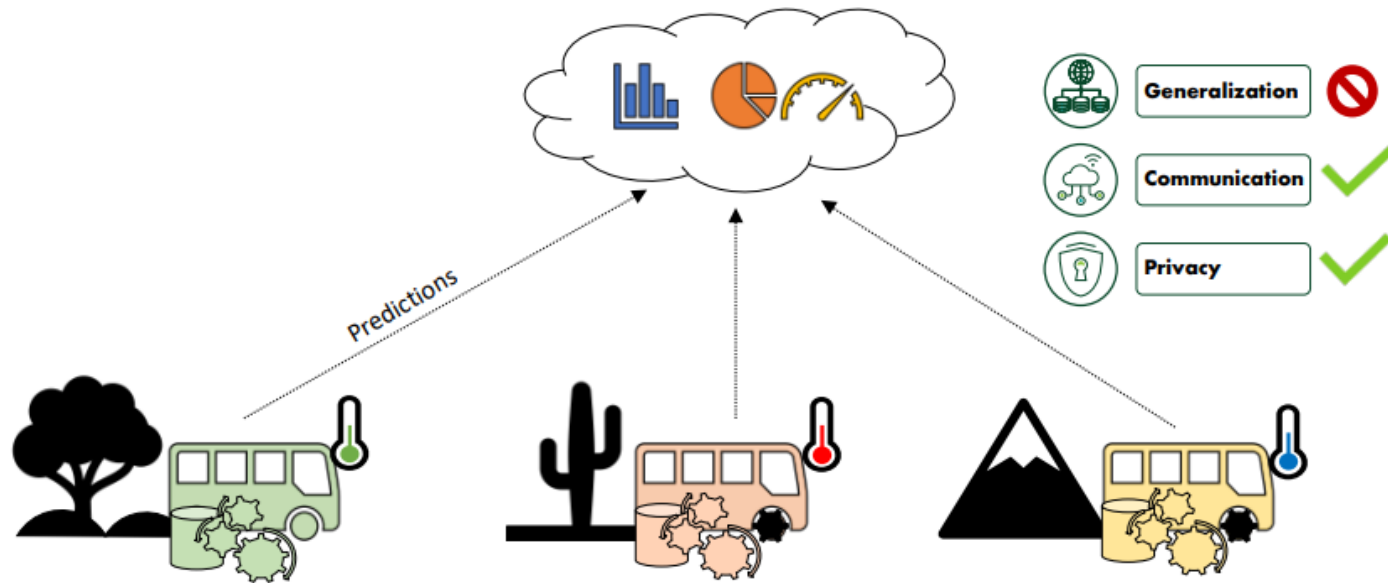
Edge/decentralized learning: ML continuously onboard each device at the edge of the network.

■ ML that runs on site, onboard each connected device, by continuously training the ML model on streaming data, the devices learn an individual model for their environment.

■ Each model only needs to be able to explain what is normal for itself and not how it varies compared to all other devices.

■ Models adapt to changes over time, learning is not constrained by the internet connection and that no confidential information needs to be transferred to the cloud.

■ It is not possible to get an overall view and learning

Learning: federated

Federated Learning – learning one from each other while keeping data on the device – creation of Super Models.

- ML technique to train algorithms across decentralized edge devices while holding data samples locally

- Google started as the main player

- Aim to train ML models on billions of mobile phones while respecting the privacy of the users

  - Only send fractions of training results, i.e. training derivatives, to the cloud
  - Never store anything on the device

Federated Learning – learning one from each other while keeping data on the device – creation of Super Models.
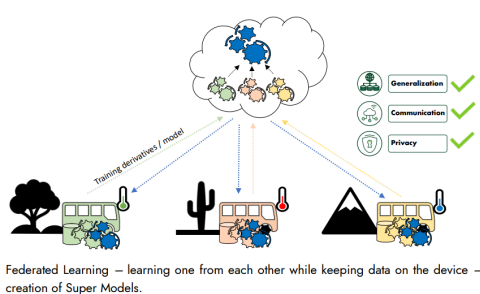
Learning: federated

- When collected in the cloud, the partial training results can be assembled to a new supermodel that, in the next step, can be sent back to the devices

  - **Goolgle open source framework TensorFlow Federated**

- Model inspection – evaluation of device behavior through its model

- Model comparison – comparing models in the cloud to find outliers, super models

- Robust learning - learning can continue even if connection to the cloud is lost

- Tailored initialization – new devices can start with a model from a similar device, instead of a general super model

Federated learning: iterative

- FL employs an iterative method containing multiple client-server exchanges: federated learning round
  - Diffuse the current/updated global model state to the contributing nodes (participants)
  - Train the local models on those nodes to yield certain potential model updates from the nodes
  - Process and aggregate the updates from local nodes into an aggregated global update so that the central model can be updated accordingly

- FL server is used for this processing and aggregation of local updates to global updates
  - Local training is performed by local nodes with respect to the commands of FL server

Learning: federated

- The FL approach allows for mass processing of data in a distributed way
- It can follow a client-server architecture

  Server sends the model to be created to the clients (1 – green lines)

  Results of the local computation are sent to the server, which aggregates them into the global model (2 – blue lines)

  Returns the new aggregated model to the clients (3 – red lines)

  This iteration, named federated learning round (FLR), occurs until some stopping criterion is reached, such as model convergence or maximum number of iterations reached.

  Edge devices only send information from their local models (parameters, hyperparameters (before training), weights, etc).
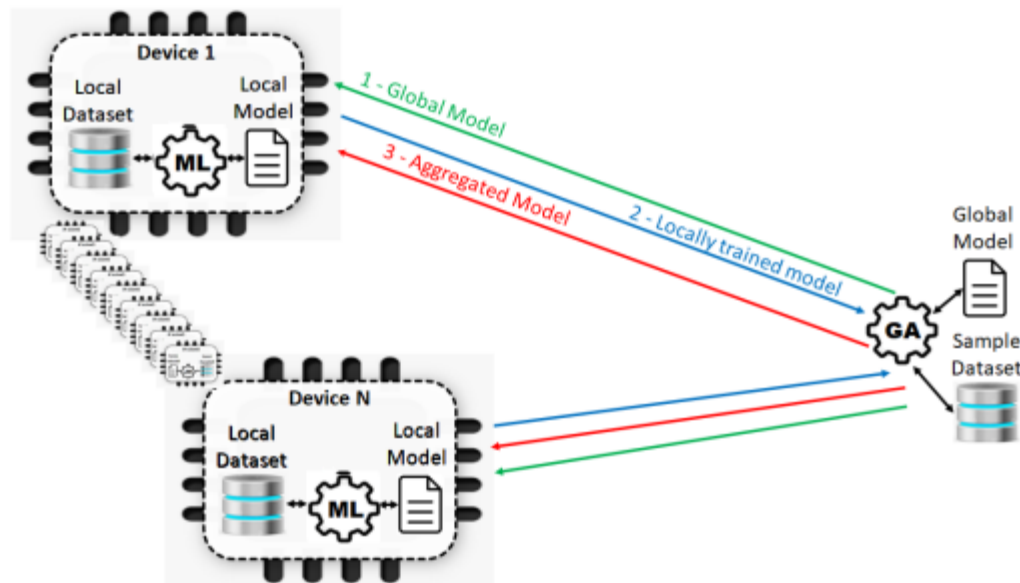
Learning: federated

- Federated learning distributes deep learning by eliminating the necessity of pooling the data into a single place

- In FL, the model is trained at different sites in numerous iterations

Learning: model aggregation

- Effective aggregation of distributed models across devices is essential for creating a generalized global model.

- Its efficiency affects precision, convergence time, number of rounds and network overhead.

- Federated stochastic gradient descent (SGD): uses a single instance of the dataset to perform the local training on the client per round of communication. SGD requires a substantial number of training rounds to produce reliable models. This algorithm is the baseline of federated learning.

- FedAvg algorithm starts from the SGD, but each client locally performs a train using the local data at the current model with multiple steps of SGD before sending the models back to the server for aggregation

FedAvg reduces the communication overhead required

to upload and download the FL model

It requires clients to perform more total computation

during training.

Local epoch: one complete pass of the training dataset through

the algorithm

Algorithm 1 Algorithm FedAvg. The K clients are indexed by k; B is the local minibatch size, E is the number of local epochs, and n is the learning rate [3].

function SERVER-SIDE:
    initialize w0
    for each round $t = 1,2,...$ **do**
        $m \leftarrow \max(C{\cdot}k,1)$
        $S_t \leftarrow$ (random set of m clients)
        for each client $k \in S_t$ **in parallel do**
            $w_{t+1} \leftarrow$ ClientUpdate(k, wt)
        end for
        $w_{t+1} \leftarrow \sum_{k=1}^{K} \frac{n_k}{n} w_{t+1}^{k}$
    end for
end function

function CLIENT-SIDE:
    ClientUpdate(k,w): // Run on client k
    $B \leftarrow$ (split P_k into batches of size B)
    for each local epoch i from 1 to E **do**
        for batch $b \in B$ **do**
            $w \leftarrow w - \eta \nabla l\,(w; b)$
        end for
    end for
    return w to server
end function

Learning: model aggregation

- Fault Tolerant Federated Average: ability of a computing system to continue working in the event of a failure

    Can tolerate some nodes being offline during secure aggregation

- Q-Federated Average: re-weight the objective in order to achieve fairness in the global model

    Gives higher weights to devices with poor performance

    The network's accuracy distribution becomes more uniform

    $$\min_w f_q(w) = \sum_{k=1}^{m} \frac{p_k}{q+1} F_k^{q+1}(w)$$

    Fk(·) to the power of (q+1), q is a parameter that tunes the amount of fairness to impose.

- Federated Optimization: uses a client optimizer during the multiple training epochs and a server optimizer during model aggregation

    ADAGRAD, ADAM, and Yogi

Algorithm 1 FEDOPT

1: Input: x_0, CLIENTOPT, SERVEROPT
2: **for** $t = 0, \cdots, T-1$ **do**
3:     Sample a subset $\mathcal{S}$ of clients
4:     $x_{i,0}^t = x_t$
5:     **for each client** $i \in \mathcal{S}$ **in parallel do**
6:         **for** $k = 0, \cdots, K-1$ **do**
7:             Compute an unbiased estimate $g_{i,k}^t$ of $\nabla F_i(x_{i,k}^t)$
8:             $x_{i,k+1}^t = \text{CLIENTOPT}(x_{i,k}^t, g_{i,k}^t, \eta_l, t)$
9:         $\Delta_i^t = x_{i,K}^t - x_t$
10:    $\Delta_t = \frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} \Delta_i^t$
11:    $x_{t+1} = \text{SERVEROPT}(x_t, -\Delta_t, \eta, t)$

Algorithm 2 FEDADAGRAD, FEDYOGI, and FEDADAM

1: Initialization: $x_0, v_{-1} \geq \tau^2$, decay parameters $\beta_1, \beta_2 \in [0, 1)$
2: **for** $t = 0, \cdots, T-1$ **do**
3:     Sample subset $\mathcal{S}$ of clients
4:     $x_{i,0}^t = x_t$
5:     **for each client** $i \in \mathcal{S}$ **in parallel do**
6:         **for** $k = 0, \cdots, K-1$ **do**
7:             Compute an unbiased estimate $g_{i,k}^t$ of $\nabla F_i(x_{i,k}^t)$
8:             $x_{i,k+1}^t = x_{i,k}^t - \eta_l g_{i,k}^t$
9:         $\Delta_i^t = x_{i,K}^t - x_t$
10:    $\Delta_t = \frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} \Delta_i^t$
11:    $m_t = \beta_1 m_{t-1} + (1 - \beta_1)\Delta_t$
12:    $v_t = v_{t-1} + \Delta_t^2$ (**FEDADAGRAD**)
13:    $v_t = v_{t-1} - (1 - \beta_2)\Delta_t^2 \text{sign}(v_{t-1} - \Delta_t^2)$ (**FEDYOGI**)
14:    $v_t = \beta_2 v_{t-1} + (1 - \beta_2)\Delta_t^2$ (**FEDADAM**)
15:    $x_{t+1} = x_t + \eta \frac{m_t}{\sqrt{v_t} + \tau}$

TensorFlow Federated

- Open source framework for experimenting with machine learning and other computations on decentralized data.
- Locally simulating decentralized computations into the hands of all TensorFlow users.
  - ML model architecture of our choice
  - Train it locally across data by all users
- Version of the NIST dataset that has been processed by the Leaf project to separate the digits written by each volunteer.

```python
1   # Load simulation data.
2   source, _ = tff.simulation.datasets.emnist.load_data()
3   def client_data(n):
4     dataset = source.create_tf_dataset_for_client(source.client_ids[n])
5     return mnist.keras_dataset_from_emnist(dataset).repeat(10).batch(20)
6
7   # Wrap a Keras model for use with TFF.
8   def model_fn():
9     return tff.learning.from_compiled_keras_model(
10        mnist.create_simple_keras_model(), sample_batch)
11
12  # Simulate a few rounds of training with the selected client devices.
13  trainer = tff.learning.build_federated_averaging_process(model_fn)
14  state = trainer.initialize()
15  for _ in range(5):
16    state, metrics = trainer.next(state, train_data)
17    print (metrics.loss)
```

federated_learning_for_image_classification.py hosted with ♥ by GitHub

TensorFlow Federated

- Training an ML model with federated learning is one example of a federated computation

- Evaluating it over decentralized data is another
  - Array of sensors capturing temperature readings
  - Compute the average temperature across these sensors

- Each client computes its local contribution

- Centralized coordinator aggregates all the contributions.

```python
1    @tff.federated_computation(READINGS_TYPE)
2    def get_average_temperature(sensor_readings):
3      return tff.federated_average(sensor_readings)
```

get_average_temperature.py hosted with ♥ by GitHub

Flower: A Friendly Federated Learning Framework (https://flower.dev/)

- Open source framework for experimenting with machine learning and other computations on decentralized data.

- Able to use in containers in a federated framework, FedFramework

```
List containers:

http://10.0.22.37:8000/containers/list


Server Creation:

http://10.0.22.37:8000/containers/create/server?img_server=fed-
server&port=5010&id=10&clients=4&algorithm=FedAvg&model=cnn&rounds=10&epochs=5
&predict=true



Rapid deployment of a testbed and run tests:

http://10.0.22.37:8000/run?img_server=fed-server&img_client=fed-
client&model=cnn&clients=4&rounds=10&epochs=5&predict=true&predict_client=true
```

```
Containers:
/containers/list
/containers/create/server
/containers/create/client
/containers/start
/containers/stop
/containers/remove
```
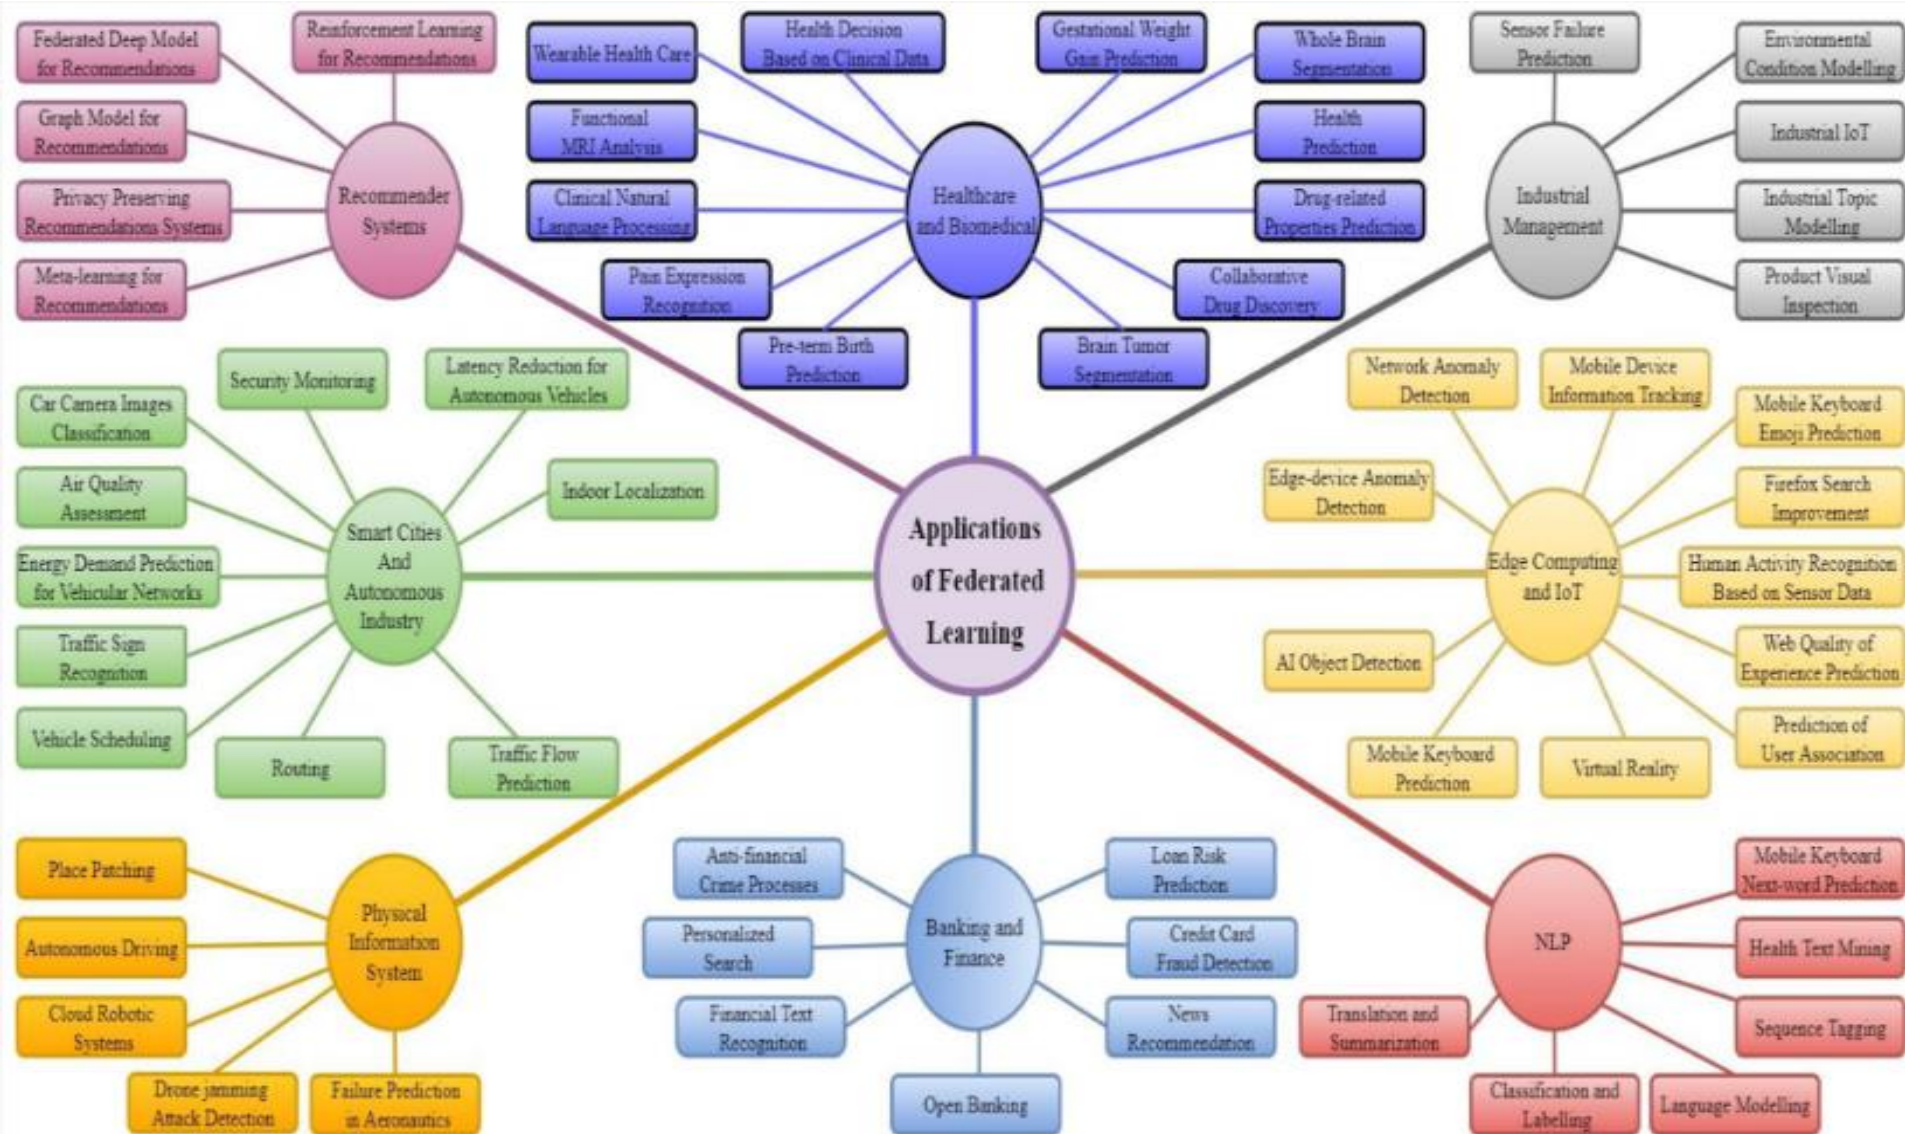
Applications for Federated

Applications of Federated Learning

Recommender Systems
- Federated Deep Model for Recommendations
- Reinforcement Learning for Recommendations
- Graph Model for Recommendations
- Privacy Preserving Recommendations Systems
- Meta-learning for Recommendations

Healthcare and Biomedical
- Wearable Health Care
- Health Decision Based on Clinical Data
- Gestational Weight Gain Prediction
- Whole Brain Segmentation
- Functional MRI Analysis
- Health Prediction
- Clinical Natural Language Processing
- Drug-related Properties Prediction
- Pain Expression Recognition
- Collaborative Drug Discovery
- Pre-term Birth Prediction
- Brain Tumor Segmentation

Industrial Management
- Sensor Failure Prediction
- Environmental Condition Modelling
- Industrial IoT
- Industrial Topic Modelling
- Product Visual Inspection

Smart Cities And Autonomous Industry
- Security Monitoring
- Latency Reduction for Autonomous Vehicles
- Car Camera Images Classification
- Indoor Localization
- Air Quality Assessment
- Energy Demand Prediction for Vehicular Networks
- Traffic Sign Recognition
- Vehicle Scheduling
- Routing
- Traffic Flow Prediction

Edge Computing and IoT
- Network Anomaly Detection
- Mobile Device Information Tracking
- Mobile Keyboard Emoji Prediction
- Edge-device Anomaly Detection
- Firefox Search Improvement
- Human Activity Recognition Based on Sensor Data
- AI Object Detection
- Web Quality of Experience Prediction
- Mobile Keyboard Prediction
- Virtual Reality
- Prediction of User Association

Physical Information System
- Place Patching
- Autonomous Driving
- Cloud Robotic Systems
- Drone jamming Attack Detection
- Failure Prediction in Aeronautics

Banking and Finance
- Anti-financial Crime Processes
- Loan Risk Prediction
- Personalized Search
- Credit Card Fraud Detection
- Financial Text Recognition
- News Recommendation
- Open Banking

NLP
- Mobile Keyboard Next-word Prediction
- Health Text Mining
- Translation and Summarization
- Sequence Tagging
- Classification and Labelling
- Language Modelling

Applications for Federated

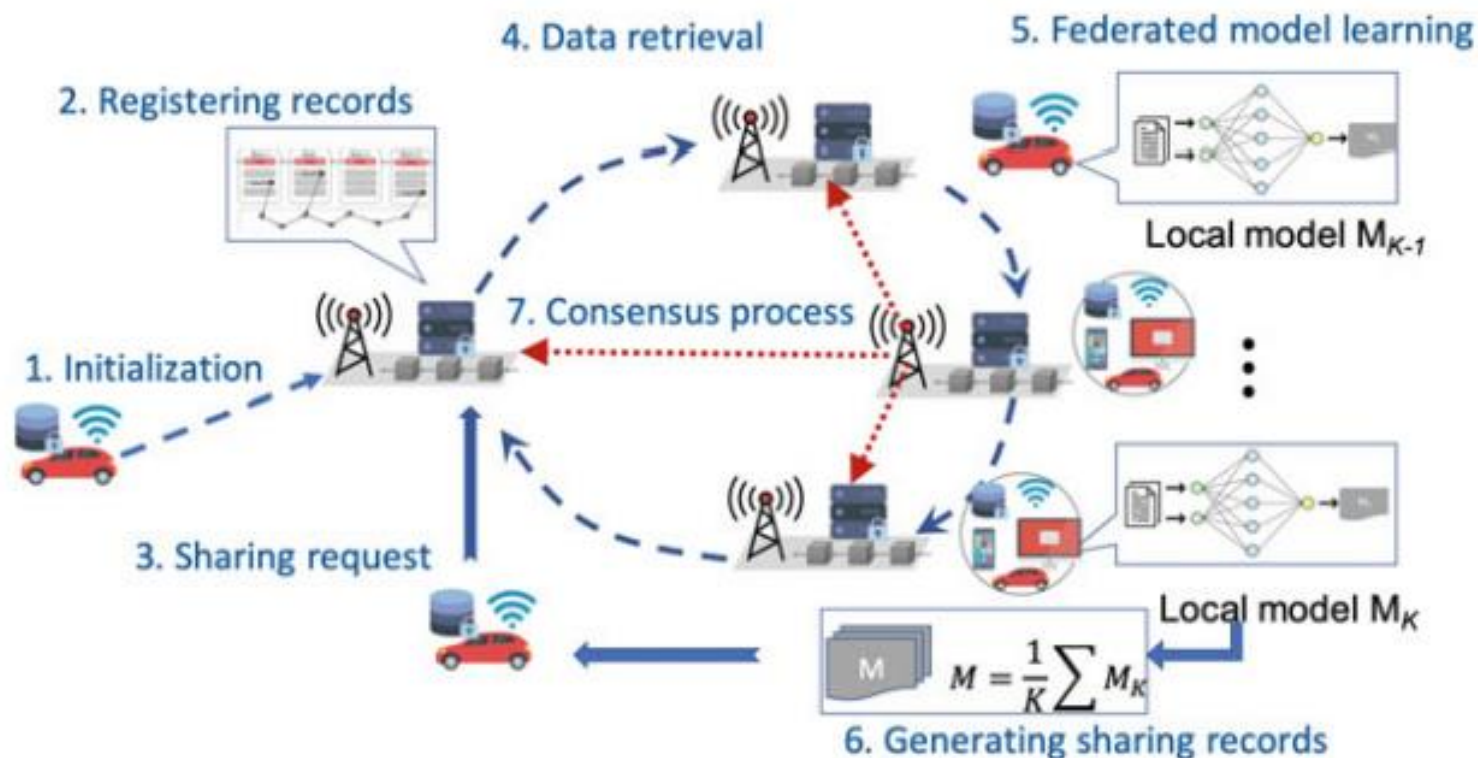| Domain	Applications
Edge computing	FL is implemented in edge systems using the MEC (mobile edge computing) and DRL (deep reinforcement learning) frameworks for anomaly and intrusion detection.
Recommender systems	To learn the matrix, federated collaborative filter methods are built utilizing a stochastic gradient approach and secured matrix factorization using federated SGD.
NLP	FL is applied in next-word prediction in mobile keyboards by adopting the FedAvg algorithm to learn CIFG [93].
IoT	FL could be one way to handle data privacy concerns while still providing a reliable learning model
Mobile service	The predicting services are based on the training data coming from edge devices of the users, such as mobile devices.
Biomedical	The volume of biomedical data is continually increasing. However, due to privacy and regulatory considerations, the capacity to evaluate these data is limited. By collectively building a global model for the prediction of brain age, the FL paradigm in the neuroimaging domain works effectively.
Healthcare	Owkin [31] and Intel [32] are researching how FL could be leveraged to protect patients' data privacy while also using the data for better diagnosis.
Autonomous industry	Another important reason to use FL is that it can potentially minimize latency. Federated learning may enable autonomous vehicles to behave more quickly and correctly, minimizing accidents and increasing safety. Furthermore, it can be used to predict traffic flow.
Banking and finance	The FL is applied in open banking and in finance for anti-financial crime processes, loan risk prediction, and the detection of financial crimes.

Federated learning in self-driving

- Edge vehicles compute the model locally; after completing each local training epoch, they retrieve the global model version and compare it to their local version.

- In order to form a global awareness of all local models, the central server performs aggregation based on the ratio determined by the global and local model versions.

- The aggregation server returns the aggregated result to the edge vehicles that request the most recent model.
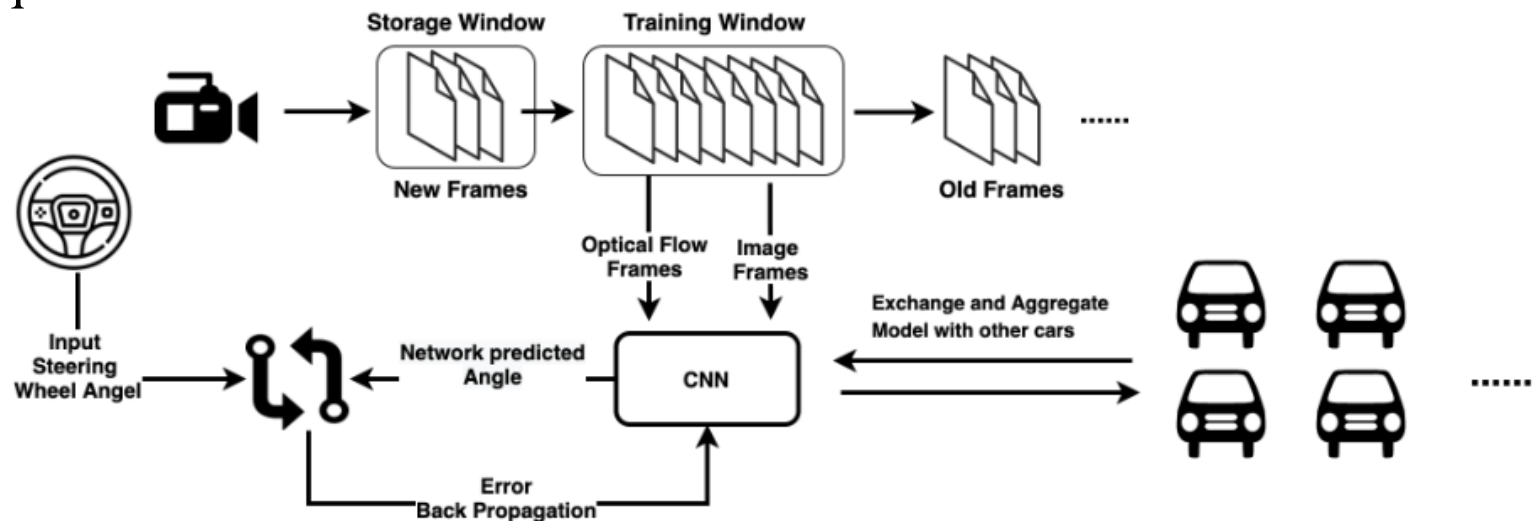
Edge-based Federated

- MEC-empowered model sharing
  - **Edge intelligence to wireless edge networks and enhances the connected intelligence among end devices in 6G networks.**

- https://www.pdl.cmu.edu/SDI/2019/slides/2019-09-05Federated%20Learning.pdf

- https://wires.onlinelibrary.wiley.com/doi/epdf/10.1002/widm.1443

- https://medium.com/tensorflow/introducing-tensorflow-federated-a4147aa20041